

1. R as a Calculator (for Scalars)

Command	Meaning	Example
Arithmetic: x [+-*/^] y x %/% y x %% y	$x + y, x - y, xy, x/y, x^y$ integer division modulo (remainder)	$7 / 3, 8^{(1/3)}$ $7 \% \% 3$ $7 \% \% 3$
Calculator functions: exp() log(x, base = exp(1)) ("=" indicates default) cos(), sin(), tan() sqrt()	exponential logarithm trigonometry square root	exp(1) log(9, base = 3) e = exp(1); log(e^2) sin(pi/2) sqrt(9)
Other easy functions: abs(x) floor(x) ceiling(x) round(x, digits = 0) signif(x, digits = 6)	absolute value greatest int $\leq x$ smallest int $\geq x$ round to #decimal places round to #significant	abs(-3) floor(-1.5) ceiling(-1.5) round(4/3, 2) signif(4/3, 2)
Statistics distributions: dnorm(x, mean = 0, sd = 1) pnorm(q, mean = 0, sd = 1) qnorm(p, mean = 0, sd = 1) rnorm(n, mean = 0, sd = 1) [dpqr] [t, chisq, f, binom] ()	$f(x)$ $P(X \leq q)$ for $X \sim N(\text{mean}, \text{sd})$ x with $P(X \leq x) = p$ random from $N(0, 1)$ other distributions	dnorm(0) # density pnorm(-1, 0, 1) # probability qnorm(.16, 0, 1) # quantile rnorm(1, 7, .01) # random ?pt, pt(-2, 100)
Miscellaneous: ?name ??topic <- (or =) variable.name ls() rm(list = ls()) list.files() # quit() demo(topic) source(file) setwd(dir)	 help("name") help.search("topic") assign variable print(variable.name) list variables clear all variables list all files comment rest of line quit R run demo code read code from file set working directory	 ?pt (help includes Description, Usage, Arguments, Value, Examples) ??deviation x <- 3 (or x = 3) x N <- 3 # number of points demo("graphics"), demo("plotmath") source("quiz1.R") setwd("C:/Users/jg/Desktop/327")
Shortcuts ... ↑, ↓ (up-, down-arrow) Esc ...	 previous command, next interrupt current command	Help > Keyboard Shortcuts

2. Vector

A *vector* (or one-dimensional *array*) `v` is a collection of *values* (or *elements*) of the same type, each identified by an *index* in the range 1 to `length(v)`. *Combine* values into a vector with `c(...)`. e.g.

```
v <- c(2.71, 5, 3.14)
```

```
length(v)
```

```
v
```

index i	value v[i]
1	2.71
	5
3	

```
words <- c("tree", "ant", "chainsaw")
```

```
length(words)
```

```
words
```

index i	value words[i]
1	
2	"ant"
	"chainsaw"

Basic Vector Types, and Specifying Constants of These Types

- **numeric** (real number): digits with optional decimal point, with optional suffix of `E` or `e` for *exponent* digits (scientific notation); e.g. `3.14e2` is _____
- **character** (which should have been called *character string*): a *string* (or word) in double or single quotes, `"..."` or `'...'`. (*Escape sequences* include `\"` (double quote), `\'` (single quote), `\n` (newline), `\t` (tab), and `\\` (backslash).)

`paste(..., sep = " ")` makes a string from its arguments, separated by `sep`. e.g.

```
oak <- 70
```

```
text = paste(sep="", "Tree names include \"oak.\\nOak weighs ", oak, " lbs/ft^3.\\n")
```

`cat(..., sep = " ")` pastes and writes to console, interpreting escape sequences. e.g.

```
cat(text)
```

```
cat(sep = "", "oak=", oak, "\\n") # display variable with helpful label
```

- **logical**: `TRUE` and `FALSE` (which become 1 and 0 when used in arithmetic)

`any(v)` is `TRUE` if any of the values in `v` is `TRUE`; `all(v)` is `TRUE` if all are

e.g. `v > 3, words == "ant", sum(v > 3), sum(words == "ant")`

`vector(mode="logical", length=0)` creates a vector of the given `mode` and `length`.

To change a vector's type, use `as.numeric()`, `as.character()`, or `as.logical()`. (There are three other basic types we will not use much: `integer`, `complex`, and `raw`.)

Names attribute

`names(x)` gets or sets a vector of **character** (strings) corresponding to values in `x`. e.g.

```
names(v) = c("e", "five", "pi"); v # set names
```

```
names(v) = NULL; v # remove names
```

Names can also be set with `c()` by specifying “name=value” pairs. e.g. `y = c(burger=2.50, fries=1.50); y`

A Few Functions

e.g. `x <- c(12, 11, 16, 11)`

`sum(x), max(x), mean(x), median(x), sd(x)`

Operators (which act element-wise on vectors)

- arithmetic: `+` `-` `*` `/` `^` (and, for integer division, `/%` is quotient, `%%` is remainder)

e.g. The sample standard deviation of x_1, x_2, \dots, x_n is $s_x = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$:

```
n = length(x)
sqrt(sum((x - mean(x))^2) / (n-1))
```

- relation: `>` `>=` `<` `<=` `==` `!=` (last two are *equals* and *is not equal to*)

• logic:	! (“not”)	T	F	& (“and”)	T	F	(“or”)	T	F
		F	T	T	T	F	T	T	T
				F	F	F	F	T	F

e.g. `!(v > 3)`, `v < 4`, `(v > 3) & (v < 4)`, `(v > 3) | (v < 4)`

- assignment: `<-` (or `=`, which is not `==`)

- sequence: `:` (colon); e.g. `11:14` is `c(, , ,)`

`seq()` is a related function:

– `seq(from=1, to=1, by)`, e.g. `seq(10, 15, by=2)` is `c(, ,)`

– `seq(from=1, to=1, length.out)`, e.g. `seq(10, 15, length.out=3)` is `c(, ,)`

- matching: `%in%`, e.g. `1:3 %in% c(2, 7)` is `c(, ,)`

Indexing

- For a vector `v` of positive integer, `x[v]` is those elements of `x` with indices in `v`; e.g. for `x <- 11:20` and `v <- c(1, 2, 10)`, `x[v]` is `c(, ,)`; `x[3]` is `c()` (or _____)

- For a vector `v` of negative integer, `x[v]` is those elements of `x` *excluding* those with indices in `v`; e.g. for `x <- 11:20` and `v <- c(-1, -2, -10)`, `x[v]` is _____

- For a vector `v` of logical,

– `which(v)` is a vector of *indices* for which `v[i]` is TRUE; e.g.

```
indices = which(x < 14) # c( , , )
```

Now use the indices: `x[indices]` is `c(, ,)`

– `x[v]` is those elements of `x` corresponding to TRUE values in `v`. e.g.

```
x[x < 14] is c( , , )
```

(so “which” could have been omitted in previous example)

e.g. `x[(x %% 2) == 0]`

- For a vector `v` of character names, `x[v]` is those elements of `x` whose names are in `v`; e.g.

```
x <- 1:3; names(x) <- c("one", "two", "Fred"); v <- c("Fred", "one"); x[v] is _____
```

3. Vector (continued) and List

Vector (continued)

Sorting functions

`sort(x, decreasing = FALSE)` returns a sorted copy of `x`. e.g. `x = c(12, 11, 16, 11); sort(x)`
`v = rank(x, ties.method = "average")`: `v[i]` is rank of `x[i]`; also try `ties.method = "first"`
`v = order(x, ..., decreasing = FALSE)`: `v[i]` is index of *i*th smallest value in `x`, so `x[v]` is sorted. (`order()` sorts data frames, coming soon. More vectors may be given in `...` to break ties.)

Structure, summary, quantile

`str(object)` displays the structure of `object`; e.g. `str(x)`
`summary(object)` summarizes it; e.g. `s = summary(x)`
`v = quantile(x, probs = seq(0, 1, 0.25))`: `v[i]` is the quantile corresponding to `probs[i]`

NULL and NA special values

NULL is the *null object* returned by expressions and functions whose value is undefined; e.g. `names(x)`

NA (“not available”) indicates a missing data value. NA propagates in calculations. e.g.

`x[3] = NA; sum(x); sum(x, na.rm=TRUE) # "na.rm=FALSE" is a common function argument`

Test for these values with `is.null()` and `is.na()`, not with “== NULL” or “== NA”

(Vector) File input/output

`scan(file = "", what = numeric())` reads a vector of `numeric` from `file`, which defaults to `stdin` (the console). `what` options include `logical()`, `numeric()`, and `character()`.

`write(x, file = "data")` writes vector `x` to `file` (which defaults to `"data"`). e.g.

`fifties = 50:59; write(x=fifties, file="50s.txt"); y = scan(file="50s.txt", what=integer())`

List

A list is an ordered collection of components not necessarily of the same type. e.g.

```
x = list("rope", 72, c(5, 7, 10), TRUE); length(x)
```

List components can be named via name=value pairs in `list()` or via `names()`.

```
y = list(self = "John",
        spouse = "Michele",
        kids.ages = c(0, 2, 5, 7, 9, 11)
      )
```

List indexing with single brackets returns a sub-list

(Recall: vector indexing with `[...]` returns a sub-vector.) The index can be a vector of **integer** (select), **negative integer** (exclude), **logical** (select TRUE), or **character** names (select). e.g.

```
str(y[2]); str(y[2:3]); str(y[c(-2, -3)]); str(y[c("spouse", "kids.ages")])
```

Indexing with double brackets or `$` returns one component, dropping names

e.g. `x[4]` vs. `x[[4]]`

e.g. `str(y[[2]])`; `str(y[["spouse"]])`

The `$` operator used in the form `list.name$component.name` is the same as `list.name[[component.name]]`, except that `$` doesn't allow a computed index.

e.g. `y$spouse`; `y$kids.ages`; `y$kids.ages[3]`

e.g. `z = "spouse"`; `y[[z]]`; `y$z`

Add a component to a list by assigning it:

```
y$kids.names = c("Teresa", "Margaret", "Monica", "Andrew", "Mary", "Philip")
```

Remove a component from a list by setting it to NULL: `y$self = NULL`

e.g. Sort names alphabetically; then sort ages to keep up:

```
indices.ordered.by.name = order(y$kids.names); indices.ordered.by.name
(names.sorted = y$kids.names[indices.ordered.by.name]) # (...) calls print(...)
(ages.sorted.by.name = y$kids.ages[indices.ordered.by.name])
```

Convert a list to a vector

`unlist(x, use.names=TRUE)` simplifies a list `x` to a vector (where possible). e.g. `unlist(y)`

A *data frame*, coming soon, is (\approx) a list of equal-length vectors (like a spreadsheet).

4. Data Frame, Factor, Formula

Data Frame (R's fundamental data structure)

A `data.frame` is (\approx) a list of equal-length vectors.

e.g. `mtcars` is a built-in `data.frame`: `mtcars`, `?mtcars`, `str(mtcars)`, `summary(mtcars)`.

Factor

A *factor* represents a vector of categorical values. Categorical data must be converted to factors for R's summary, plotting, and modeling functions to work correctly.

`factor(x, levels, labels = levels)` makes a factor from vector `x`, using levels in `levels` (or the unique strings in `as.character(x)` by default), using optional `labels` to make the categories more readable. e.g.

```
m = mtcars          # work on a copy
m$vs = factor(mtcars$vs, labels=c("V", "straight"))
m$am = factor(mtcars$am, levels=c("0", "1", "2"), labels=c("auto", "man", "CVT"))
# suppose level "2" and label "CVT" are useful even though they are not in mtcars
str(m)
summary(m) # now categorical variables are handled well
```

`table(...)` makes a contingency table of counts of each combination of factors in e.g.
`table(m$vs)`, `table(m$vs, m$am)`

Data frame manipulation examples

```
m$mpg          # mpg column
m[, 1]         # all rows, 1st column (mpg again)
m[1:3, 1:3]    # rows 1:3, columns 1:3
dim(m)         # dimensions
n.rows = dim(m)[1]
n.cols = length(m) # or dim(m)[2]
tail(m)
rownames(m)[n.rows - 2] = "Monica's present"
m$hp[30] = 25
M = median(m$hp)
mean(m$mpg[m$hp > M]) # high-power mileage
mean(m$mpg[m$hp < M]) # low-power mileage
m$price = 1000*(1:n.rows) # add column
m$vs = NULL          # delete column
sorted = m[order(m$cyl, m$disp), ] # sort by cyl, then by disp
```

(Data frame) File input and output (and “.csv” for Excel)

- `write.table(x, file = "", ...)` writes `x` to file. Variants include `write.csv(x, file = "")`.
e.g. `write.csv(m, file = "mtcarsMonica.csv")` saves `m` (our corrupted `mtcars`) as comma-separated values (csv)
- `table = read.table(...)` reads from a file into a `data.frame`. Variants include
 - `table = read.csv(file, header = TRUE, row.names = 1)` for a file of comma-separated values with a header row of column names and a first column of row names; e.g.
`monica = read.csv("mtcarsMonica.csv", row.names = 1)`
 - `table = read.csv(file, header = FALSE, col.names = c(...), row.names = c(...))` for a file of unlabeled data

Formula

A *formula* of the form `y ~ model` indicates that `y` depends on the variables in `model`. e.g. Here's a preview of the use of formulas in the coming handouts on graphics and regression.

- Here's a lousy boxplot that obscures the dependence of flower length on flower species:

```
flowers = read.csv("flowers.csv")
str(flowers) # note the factor
boxplot(flowers$Flower.Length,
        main="Flower Length Without Regard for Species", ylab="Length (mm)")
```

Improve the graph: `boxplot(formula, ...)` makes multiple plots of data specified by `formula`. e.g. This triple plot reveals the dependence of length on species as a grouping variable:

```
boxplot(flowers$Flower.Length ~ flowers$Species,
        main="Flower Length by Species", xlab = "Species", ylab="Length (mm)")
```

- Here are similar examples using `mtcars`:

```
boxplot(m$disp)
boxplot(m$disp ~ m$am)
```

- We'll use formulas in linear regression soon:
 - `y ~ x` indicates that y depends linearly on x , as in the simple linear regression model,
 $y = a_1 + a_2x$
 - `y ~ x1 + x2 + x3 + x1*x2` indicates that y depends linearly on x_1, x_2, x_3 , and $x_1 \cdot x_2$,
as in the multiple linear regression model, $y = a_1 + a_2x_1 + a_3x_2 + a_4x_3 + a_5x_1 \cdot x_2$.

5. (Base) Graphics

Common parameters

- `formula`, `data`: a formula of the form `y ~ model` and a data frame containing the variables
- `main`, `sub`; `xlab`, `ylab`: main title, subtitle; *x*-axis, *y*-axis labels
- `xlim`, `ylim`, each a 2-vector (low, high): *x*-axis, *y*-axis limits
- `pch`: plotting character (see `?points`)
- `cex` (symbols), `cex.axis`, `cex.lab`, `cex.main`, `cex.sub`: character expansion (relative to 1)
- see `?par` for others

Numeric data

- `boxplot(x)` makes a boxplot from vector `x`; `boxplot(x ~ g)` groups by factor `g`; e.g.

```
boxplot(mtcars$mpg, main="Gas mileage", ylab="miles per gallon", ylim=c(0,40))  
boxplot(mpg ~ factor(cyl), data=mtcars, xlab="cylinders", ylab="miles per gallon")
```
- `stripchart(x, method="overplot")` makes a dot plot of `x` (better than boxplot for small sample); `stripchart(x ~ g)` groups by `g`; `method` handles duplicates: "overplot", "jitter", or "stack"; e.g. `stripchart(mpg ~ factor(am), data=mtcars, method="stack")`
- `hist(x, breaks="Sturges", freq=NULL)`, makes a histogram from `x`, where `breaks` is a vector of bin boundaries (or, as in the default "Sturges", the name of a bin algorithm); `freq=FALSE` gives density histogram instead of frequency; e.g. `hist(mtcars$mpg)`
- `plot(x, y)` makes a scatterplot from vectors `x` and `y`; e.g. `x = 1:5; y = 2*x; plot(x, y), plot(x, y, xlim=c(0,10), ylim=c(0,10))`
- `points(x, y)` adds points to a plot, and `lines(x, y)` adds line segments; e.g. `points(x, x, pch=15); lines(x=c(1,3,5,7,9), y=c(8,1,4,1,8), col="red")`
- `plot(density(x))` makes a *density plot* (usually better than a histogram) from `x`; `rug(x)` adds the data points; e.g. `plot(density(mtcars$mpg)); rug(mtcars$mpg)`
(note: `density(x)` estimates density $f(x)$, returning a list including `(x, y)`, where $y \approx f(x)$)
- `pairs(x)` makes a matrix of scatterplots of pairs of columns of data frame `x`; e.g. `pairs(mtcars)`
- `curve(expr, from=NULL, to=NULL, n=101, add=FALSE, type="l")` draws a curve of `expr` over `[from, to]` (`add=TRUE` \implies add to existing plot); e.g.

```
curve(expr=x*sin(1/x), from=-pi/6, to=pi/6, n=200); curve(expr=x*1, add=TRUE, col="red")
```


Legends; math expressions in titles and labels

`legend(x, y, legend, col=par("col"), lty, pch)` makes a legend at (x, y) (or x can be one of {"bottomright", etc.}: see `?legend`) using labels, colors, line types, and plotting characters in vectors `legend`, `col`, `lty`, and `pch`; e.g.

```
legend("top", legend=c("x*sin(1/x)", "x"), col=c("black", "red"), lty=c(1, 1))
```

Use `expression(...)` in character string used as main or `xlab` or `ylab`; see `?plotmath`. e.g.

```
legend("top", legend=c(expression(x*sin(frac(1,x))), "x"), col=c("black", "red"), lty=c(1, 1))
```

Categorical data

- `barplot(height, names.arg = NULL)` makes a barplot of the counts in `height`, with (optional) bar labels in `names.arg`; e.g.

```
counts = table(mtcars$cyl); barplot(counts)
```

- `mosaicplot(x)` makes a mosaic plot from a table of counts from `table()`; e.g.

```
counts = table(mtcars$cyl, mtcars$gear); mosaicplot(counts)
```

Multiple figures

`matrix(data, nrow, ncol, byrow=FALSE)` fills an `nrow` \times `ncol` matrix by column from `data`
`layout(mat)`, for matrix `mat`, divides graph so *i*th figure is drawn where `mat==i` (0 \implies blank)
`layout.show(n=1)` shows outlines of next *n* figures; e.g.

```
m = matrix(data=c(1, 0, 2, 3, 3, 3), nrow=2, ncol=3, byrow=TRUE)
layout(m)
layout.show(3)
hist(mtcars$mpg)           # 1st plot: (frequency) histogram alone
plot(density(mtcars$mpg))  # 2nd plot: density plot alone
hist(mtcars$mpg, freq=FALSE) # 3rd plot: density histogram
lines(density(mtcars$mpg))  # add density plot to (3rd plot) histogram
layout(matrix(data=1, nrow=1, ncol=1)) # reset graphics device
```

Write graphical output to a file

- Open a graphical output file with, e.g., `pdf("file.pdf")`, `png("file.png")`, `jpeg("file.jpg")`, `bmp("file.bmp")`, `postscript("file.ps")`, `tiff("file.tif")`
- Make graph
- Close the file with `dev.off()`

6. Statistical Tests and Confidence Intervals

One Mean or the Difference of Two Means

`out = t.test(x, y = NULL, alternative = "two.sided", mu = 0, conf.level = .95)` tests $H_0 : \mu_X = \mu_0 = \text{mu}$ for a sample `x` from a normal population; or, if `y` is given, $H_0 : \mu_X - \mu_Y = \mu_0 = \text{mu}$, for samples `x` and `y` from normal populations. `out` is a list containing (among other things):

- `$parameter`: degrees of freedom ($n - 1$, where $n = \text{length}(x)$, if `y == NULL`; or a mess)
- `$statistic`: Student's t test statistic, $t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$; or $t = \frac{(\bar{x} - \bar{y}) - \mu_0}{\sqrt{\frac{s_X^2}{n_X} + \frac{s_Y^2}{n_Y}}}$
- `$p.value`: probability of a value at least as extreme as t under H_0
- `$conf.int`: confidence interval for μ_X (or $\mu_X - \mu_Y$) corresponding to H_1 in `alternative`
- `$estimate`: \bar{x} (or \bar{x} and \bar{y})

Other alternative choices are "less" and "greater". e.g.

```
x = rnorm(n = 10, mean = 0, sd = 1); (out = t.test(x))
x = rnorm(10, 0, 1); (out = t.test(x, mu = 2)) # rnorm() isn't part of the test!
x = rnorm(10, 0, 1); y = rnorm(10, 2, 1); (out = t.test(x, y))
x = rnorm(10, 0, 1); y = rnorm(10, 2, 1); (out = t.test(x, y, mu = -2))
```

F Test for Equality of Variances

`out = var.test(x, y, ratio = 1, alternative = "two.sided", conf.level = .95)` tests $H_0 : \frac{\sigma_X^2}{\sigma_Y^2} = \text{ratio}$ for two samples `x` and `y` from normal populations. `out` is a list containing:

- `$parameter`: degrees of freedom ($n_X - 1$ and $n_Y - 1$, where $n_X = \text{length}(x)$ and $n_Y = \text{length}(y)$)
- `$statistic`: F test statistic, $f = \frac{s_X^2/\sigma_X^2}{s_Y^2/\sigma_Y^2} = \frac{s_X^2}{s_Y^2} \cdot \frac{1}{\text{ratio}}$
- `$p.value`: probability of a value at least as extreme as f under H_0
- `$conf.int`: confidence interval for $\frac{\sigma_X^2}{\sigma_Y^2}$
- `$estimate`: $\frac{s_X^2}{s_Y^2}$

e.g. `x = rnorm(100, 0, 1); y = rnorm(10, 0, 2); (out = var.test(x, y, ratio = 1))`

e.g. `x = rnorm(100, 0, 1); y = rnorm(10, 0, 2); (out = var.test(x, y, ratio = .25))`

Chi-Squared Tests

- Goodness-of-fit:

`counts = c(...); probs = c(...); (out = chisq.test(x = counts, p = probs))`
tests H_0 : “counts came from a distribution with probabilities `probs`”. e.g.

`counts=c(12,15,17,6); probs=c(.20,.25,.40,.15); (out=chisq.test(x=counts, p=probs))`
`out` is a list containing (among other things):

- `$parameter`: degrees of freedom (`#categories - 1 == length(x) - 1`)
- `$statistic`: χ^2 test statistic for goodness-of-fit of observed `counts` to proposed `probs`
- `$p.value`: probability of a value at least as extreme as χ^2 under H_0

- Independence / Homogeneity

e.g. Consider the counts in this contingency table:

Education	Smoking status			
	Nonsmoker	Former	Moderate	Heavy
Primary	56	54	41	36
Secondary	37	43	27	32
University	53	28	36	16

To get data into the test, we need `x = matrix(data, nrow, ncol, byrow = FALSE)`, which fills an `nrow` by `ncol` matrix `x`, by column, from the vector `data`. Note that `x[,c]` is the c^{th} column of `x`, and `x[r,]` is the r^{th} row. e.g.

`(x = matrix(data = c(56,37,53, 54,43,28, 41,27,36, 36,32,16), nrow=3, ncol=4))`

The χ^2 test `out = chisq.test(x)` is for H_0 : “row and column variables are independent” (or H_0 : “the column populations have the same distribution with respect to the row variable”).

`out` is a list containing (among other things):

- `$parameter`: degrees of freedom, $(\text{\#rows} - 1) \times (\text{\#columns} - 1)$
- `$statistic`: χ^2 test statistic for independence of row and column variables (or for homogeneity of column populations with respect to row variable)
- `$p.value`: probability of a value at least as extreme as χ^2 under H_0
- `$expected`: expected counts under H_0

To use `chisq.test()` on variables in a data frame, recall that `table(...)` makes a contingency table of counts of each combination of factors in e.g.

```
table(mtcars$cyl)
table(mtcars$cyl, mtcars$gear)
```

One Proportion or the Difference of Two Proportions

- For integers x and n , `out = prop.test(x, n, p, alternative = "two.sided", conf.level = .95)` tests $H_0 : p = p_0 = p$ for a sample containing x successes in n trials. e.g.

`x = 800; n = 1000; p0 = .77; (out = prop.test(x, n, p0, correct=FALSE))`

(`correct=FALSE` disables a good continuity correction that would add to the explanation.)

`out` is a list containing (among other things):

- `$parameter`: degrees of freedom ($\# \text{categories} - 1 = 2$ (success and failure) $- 1 = 1$)
- `$statistic`: χ^2 test statistic for goodness-of-fit of observed counts, x successes (800) and $n-x$ failures ($1000 - 800 = 200$), to the distribution with expected counts, $n \cdot p$ successes ($770 = .77 \times 1000$) and $n \cdot (1-p)$ failures ($230 = (1 - .77) \times 1000$).
- `$p.value`: probability of a value at least as extreme as χ^2 under H_0
- `$conf.int`: confidence interval for p
- `$estimate`: $\hat{p} = x/n$

(I teach an equivalent z -test for this one-proportion test:

`phat = x/n; z = (phat - p0) / sqrt(p0*(1-p0)/n); (p.value = 2*pnorm(-abs(z)))`

Then `z^2` matches `out$statistic` above, and the P-values are the same.

)

- For 2-vectors x and n , `out = prop.test(x, n, alternative = "two.sided", conf.level = .95)` tests $H_0 : p_1 - p_2 = 0$ (or $p_1 = p_2$) for samples from two populations containing $x[1]$ successes in $n[1]$ trials and $x[2]$ successes in $n[2]$ trials, respectively. e.g.

`x = c(40, 87); n = c(244, 245); (out = prop.test(x, n, correct=FALSE))`

`out` is a list containing (among other things):

- `$parameter`: degrees of freedom, 4 (counts) $- 2$ (constraints due to the sample sizes) $- 1$ (parameter estimated, $\hat{p}) = 1$

Outcome	Sample	
	1	2
success	40	87
failure	$244 - 40$	$245 - 87$

- `$statistic`: χ^2 test statistic for goodness-of-fit of observed counts, $x[1]$ and $x[2]$ successes (40 and 87) and $n[1]-x[1]$ and $n[2]-x[2]$ failures ($244 - 40$ and $245 - 87$) to the distribution with corresponding expected counts based on $\hat{p} = \frac{x_1+x_2}{n_1+n_2}$
- `$p.value`: probability of a value at least as extreme as χ^2 under H_0
- `$conf.int`: confidence interval for the difference in proportions $p_1 - p_2$
- `$estimate`: a 2-vector containing $\hat{p}_1 = x[1]/n[1]$ and $\hat{p}_2 = x[2]/n[2]$

(Here, too, I teach an equivalent z -test, with `z^2 == out$statistic`, and the same P-value.)

7. Simple Linear Regression, $y = mx + b$

e.g. `cars` is a built-in data.frame: `cars`, `?cars`, `str(cars)`, `head(cars)`

- (Recall) `plot(x, y)` makes a (base graphics) scatterplot of data in the vectors `x` and `y`; e.g.
`plot(x=cars$speed, y=cars$dist)`
- `cor(x, y)` gives the correlation of vectors `x` and `y`; e.g. `r = cor(x = cars$speed, y = cars$dist)`
For data frame `x`, `A = cor(x)` gives a matrix `A` such that `A[i, j] == cor(x[, i], x[, j])`, the correlation of `A`'s i^{th} and j^{th} columns; e.g. `cor(mtcars[, 1:3])`
- `lm(y ~ x, data)` calculates a linear regression model $y = mx + b$ from the `y` and `x` variables in the data.frame `data` (this uses the “formula, data” interface mentioned earlier); e.g.
`m = lm(dist ~ speed, data = cars) # "m" is for "model"`
`str(m)`
`summary(m) # summary`
`anova(m) # ANOVA table`
- `m$coefficients` is a vector containing y -intercept b and slope m :
`y.intercept = m$coefficients[1]`
`slope = m$coefficients[2]`
- `abline(a, b)` adds a line $y = a + bx$, and `abline(reg)` adds the line from model `reg`; e.g.
`abline(a = y.intercept, b = slope) # add regression line`
`abline(reg = m) # same as previous line`
`abline(a = mean(cars$dist), b = 0, lty = "dashed") # horizontal line through mean y`
- `predict(model, newdata)` gives \hat{y} from model evaluated at x (or at x_1, \dots, x_p in the multiple regression case) in data.frame `newdata`; e.g. Our model's x is `speed`; so put speeds for which we want predictions in a data.frame with a `speed` column:

`d = data.frame(speed = seq(from=5, to = 25, by = 5))`
`y.hat = predict(m, newdata = d)`
`# add (x, y) pairs to graph with plotting character 19, scaled by 3`
`points(x=d$speed, y=y.hat, pch=19, cex=3)`
- In the simple regression model $y_i = mx_i + b + \varepsilon_i$, errors ε_i are assumed to be random and independent, with $\varepsilon_i \sim N(0, \sigma)$. To check these assumptions, a *residual plot* of points $\{(\text{fitted value} = \hat{y}_i, \text{residual} = e_i = y_i - \hat{y}_i)\}$ should show no pattern (if errors are random and independent) or varying vertical spread (if errors have the same standard deviation σ); e.g.

`plot(m$fitted.values, m$residuals)`
`abline(0, 0) # y = 0 + 0x; errors should have mean 0`

- A *QQ plot* shows quantiles of a data distribution, like our residuals, on the y -axis against the same quantiles of a reference distribution, like $N(\mu = \text{mean}(\text{residuals}), \sigma = \text{sd}(\text{residuals}))$. If the assumption of normal errors is met, these points should be close to a line. `qqline(x)` adds a line through the first and third quantile pairs. e.g.

```
x = rnorm(n=100); qqnorm(x); qqline(x) # 100 random N(0, 1) points
w = rexp(100); qqnorm(w, ylim=c(-1, 5)); qqline(w) # 100 random Exp(1) points

qqnorm(m$residuals); qqline(m$residuals) # our "dist vs. speed" model
```

Or use `plot(m)` to see the residual and QQ plots, and two others, in one step:

```
layout(matrix(data=1:4, nrow=2, ncol=2, byrow=TRUE))
plot(m)
layout(matrix(data=1, nrow=1, ncol=1)) # reset graphics device
```

Multiple Linear Regression, $y = a_0 + a_1x_1 + \dots + a_px_p$

e.g. `y ~ x1 + x2 + x3 + x1*x2` indicates that y depends linearly on x_1 , x_2 , x_3 , and $x_1 \cdot x_2$, as in the multiple linear regression model, $y = a_0 + a_1x_1 + a_2x_2 + a_3x_3 + a_4x_1 \cdot x_2$.

```
n = 100 # simulate n points, (y, x1, x2, x3), for a "sanity check" example
x1 = rnorm(n=n, mean=0, sd=1); x2 = rnorm(n); x3 = rnorm(n)
y = 3 + 4*x1 + 5*x2 + 6*x3 + 7*x1*x2
m = lm(y ~ x1 + x2 + x3 + x1*x2) # use lm() to discover coefficients from data
summary(m)

y = 3 + 4*x1 + 5*x2 + 6*x3 + 7*x1*x2 + rnorm(n) # add noise to make it harder
m2 = lm(y ~ x1 + x2 + x3 + x1*x2)
summary(m2)

m3 = lm(mpg ~ hp + wt + gear, data=mtcars) # real data from mtcars:
summary(m3)
anova(m3)
```

Inference on the coefficients is facilitated by `summary(model)`, which gives

- estimated coefficients a_0, a_1, \dots, a_p
- estimated standard deviations of coefficients, s_{a_0}, \dots, s_{a_p}
- the F statistic and P-value for $H_0 : a_1 = \dots = a_p = 0$
- for each coefficient a_i , the t statistic and P-value for $H_0 : a_i = 0$

`confint(m, level = .95)` gives confidence intervals for the coefficients

8. Simulation by replicating a calculation

Make random number generation repeatable

`set.seed(seed)`, for integer `seed`, sets starting point of (pseudo-)random number generation. e.g.

```
a = rnorm(1); b = rnorm(1); a == b
set.seed(0); a = rnorm(1); set.seed(0); b = rnorm(1); a == b
```

Repeat a calculation n times

`replicate(n, expr)` returns a vector (or matrix or array) of n evaluations of `expr`. e.g.

```
x = replicate(n=4, expr=rnorm(1))      # 4 random samples of size 1
y = replicate(n=4, expr=rnorm(3))      # 4 random samples of size 3
z = replicate(n=4, expr=mean(rnorm(3))) # 4 means of samples of size 3
```

`expr` can be compound in curly braces, `{ ... }`; its value is that of its last expression. e.g.

```
w = replicate(n=4, expr={ mu=7; sigma=3; x=rnorm(n=3, mean=mu, sd=sigma); mean(x) })
```

Distributions

Check `?distributions`. (Recall prefixes `d`, `p`, `q`, `r` for *density*, *probability*, *quantile*, *random*.)

Let's simulate a few distributions.

- $N(\mu, \sigma)$: First, confirm that \bar{x} is close to μ and that s is close to σ :

```
mu = 7; sigma = 3; mean(x <- rnorm(n=1000, mean=mu, sd=sigma)); sd(x)
```

Second, the *Central Limit Theorem* (CLT) says that for a large sample from (almost) any distribution with finite μ and σ , $\bar{X} \approx N(\mu, \frac{\sigma}{\sqrt{n}})$.

e.g. Consider $U(0,1)$, which has $\mu = \frac{\max - \min}{2} = \frac{1}{2}$ and $\sigma = \sqrt{\frac{(\max - \min)^2}{12}} = \sqrt{\frac{1}{12}}$. Simulate CLT by finding many sample means from samples from $U(0,1)$:

```
curve(dunif(x, min=0, max=1), from=-0.1, to=1.1, ylim=c(0,8), lty=2) # U(0,1)
n = 30 # sample size (also try n=1 to see CLT fail)
N = 100 # number of samples
x.bars = replicate(n=N, expr=mean(runif(n=n, min=0, max=1))) # vector of sample means
mean(x.bars) # should be near 1/2
sd(x.bars) # should be near sqrt(1/12)/sqrt(n), about .0527
curve(dnorm(x, mean=1/2, sd=sqrt(1/12)/sqrt(n)), from=0, to=1, lty=3, add=TRUE) # CLT
lines(density(x.bars), lty=1) # sampling distribution of bar(x)
rug(x.bars)
legend(x="topright", legend=c("U(0,1)", "CLT", expression(bar(X))), lty=c(2,3,1))
```

- t_{n-1} : For a random sample X_1, \dots, X_n from $N(\mu, \sigma)$, the quantity $T = \frac{\bar{X} - \mu}{s/\sqrt{n}}$ follows the *Student's t* distribution with $n - 1$ degrees of freedom, denoted t_{n-1} .

e.g. Simulate t_{n-1} for $n = 6$:

```
n = 6 # sample size
N = 100 # number of samples
mu = 7
sigma = 3
t = replicate(N, { x=rnorm(n, mean=mu, sd=sigma); (mean(x) - mu)/(sd(x)/sqrt(n)) })
plot(density(t)) # sampling distribution of T ~ t_{n-1}
rug(t)
curve(dt(x, df=n-1), lty="dashed", add=TRUE) # true t_{n-1}
curve(dnorm(x, mean=0, sd=1), lty="dotted", add=TRUE) # add N(0, 1) for reference
legend(x="topright", legend=c(expression("true " * t[n-1]), "simulated t", "N(0, 1)"),
      lty=c("dashed", "solid", "dotted"))
```

- χ_n^2 : If Z_1, \dots, Z_n are independent, $N(0, 1)$ random variables, then $X^2 = \sum_{i=1}^n Z_i^2 \sim \chi_n^2$
- F_{n_1, n_2} : If $X \sim \chi_{n_1}^2$ and $Y \sim \chi_{n_2}^2$ are independent, then $\frac{X/n_1}{Y/n_2} \sim F(n_1, n_2)$

What is a *P*-value?

A *P*-value is the probability, assuming H_0 is true, of getting data, as summarized by the test statistic, more extreme than the sample data. e.g. Guinness says pouring a glass should take 119.5 seconds. Here's a random sample of times from a server:

```
x = c(118, 121, 113, 116, 117, 112, 113)
```

Is this server pouring correctly? Test $H_0 : \mu = 119.5$ vs. $H_1 : \mu \neq 119.5$: `(out = t.test(x, mu=119.5))`

Simulate *P*-value by seeing how often `t`, from random samples, is greater than `out$statistic`:

```
mu = 119.5
sigma = sd(x)
n = length(x) # sample size
N = 1000 # number of replicates
t = replicate(N, { x=rnorm(n, mean=mu, sd=sigma); (mean(x) - mu)/(sd(x)/sqrt(n)) })
more.extreme = (abs(t) > abs(out$statistic))
(simulated.p.value = sum(more.extreme) / N)
out$p.value

plot(density(t), main=bquote(. (N) * " Simulated t statistics")) # visualize P-value
rug(t)
points(x=out$statistic, y=0, pch=19, col="red")
text(x=out$statistic, y=.02, labels="out$statistic")
```


STAT 327-1 (also -4 and -7): Introductory Data Analysis with R

Course outcome: Students will use R to manipulate data and perform exploratory data analysis using introductory statistics.

Unit	Objectives Students will:	Assessment	Read, View, Do
1 Build basic R vocabulary	<ol style="list-style-type: none"> 1. Use R as a calculator. 2. Use R's distribution functions: calculate probability mass/density, cumulative probability distribution, and quantile functions and generate random numbers. 3. Run a line of R code in the console and a batch from a script. 4. Use R Markdown to write reports integrating text, data, R code, and its output. 	Q1 HW1 (trivial script) (Q = online quiz, HW = homework)	1calculator.pdf, lecture 1
2 Manipulate data in R	<ol style="list-style-type: none"> 1. Manipulate data to create vectors, lists, and data frames. 2. Summarize a data set. 3. Select a subset of a data set. 4. Use a factor vector for categorical data. 5. Load clean tabular data sets into R. 6. Save R data sets as text or csv files. 	Q2, Q3, Q4 HW2 (donations to 2012 elections), HW3 (Boston housing)	2vector.pdf, lecture 2 3list.pdf, lecture 3 4dataFrame.pdf, lecture 4
3 Produce graphics	<ol style="list-style-type: none"> 1. Use the graphics base package to create displays of data: scatterplots, boxplots, histograms and density plots. 2. Customize graphical layout, annotations, and legends. 	Q5 HW3, HW4	5graphics.pdf, lecture 5 group work (graphics)
4 Apply statistical methods	<ol style="list-style-type: none"> 1. Run classical statistical procedures including confidence intervals, t tests, Z tests for proportions, F tests for variance, and χ^2 tests. 2. Do basic linear regression analysis and ANOVA. 	Q6, Q7 HW4 (test, regression, confidence bands)	6test.pdf, lecture 6 7regression.pdf, lecture 7 8simulation.pdf, lecture 8 group work (tests)
5 Run basic simulations	<ol style="list-style-type: none"> 1. Replicate a calculation to simulate properties of distributions. 2. Simulate data fitting $N(0, 1)$, t, χ^2, and F distributions. 		8simulation.pdf, lecture 8
		Written exam	

Prerequisite: Introductory statistics

R Markdown

R Markdown is software included with RStudio that allows you to put text, data, R code, and Latex math notation in the same plain-text file, and then compile it to a nicely formatted file containing text, data, R code, textual output of R code, graphical output of R code, and math notation. By putting all these things in a single file, R Markdown greatly simplifies the otherwise tedious and error-prone process of writing and assembling a statistical report.

Here's all you have to know for STAT 327

- To open a new file, use RStudio's menu "File > New file > R Markdown ...", give your file a name ending ".Rmd", choose "HTML" under "Default Output Format:", and click "OK".
- Write R code inside "code chunks" delimited as follows:

```
```${r}  
 # R code
```
```

(These "backquotes" are on the upper-left corner of the keyboard.)

- Write plain text anywhere in the file except in code chunks.
- Click "KnitHTML" to knit together your text, data, R code, and its output into a web page.
- For debugging, run a line of code in the console with "Ctrl-Enter" (Windows) or "Command-Enter" (Mac). See the "Chunks" menu for running a chunk at a time.

To learn more about R Markdown

Use RStudio's "?" menu to choose "Using R Markdown" and "Markdown Quick Reference".

See cheatsheets at <http://www.rstudio.com/resources/cheatsheets>. There are four; start with "R Markdown Cheat Sheet" and "R Markdown Reference Guide".

Latex for mathematical notation (optional)

In R Markdown text, you may use Latex mathematical notation in sections delimited by `$... $` to show up inline, or by `$$... $$` to show up as a separate paragraph. Here are basics:

| Latex | Result |
|----------------------------------|-----------------------|
| <code>x^y</code> | x^y |
| <code>x_y</code> | x_y |
| <code>\alpha, \mu, \sigma</code> | α, μ, σ |
| <code>\bar{x}</code> | \bar{x} |
| <code>\hat{x}</code> | \hat{x} |
| <code>\sqrt{x}</code> | \sqrt{x} |
| <code>\sum</code> | \sum |
| <code>\frac{x}{y}</code> | $\frac{x}{y}$ |

e.g. `$Z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}$` gives $Z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}$.

e.g. `$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$` gives

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i.$$

To learn more about Latex, see <http://en.wikibooks.org/wiki/LaTeX>.